

New Directions In Automated Traffic Analysis

github.com/nprint/

Jordan Holland, Paul Schmitt, Nick Feamster, Prateek Mittal

Network Traffic Analysis?

- Can we identify remote devices by probing them?
- How can we identify and stop attacks?
- Can we improve performance by analyzing traffic?
- Can users be tracked via their network traffic?

Network Traffic Analysis?


- Can we identify remote devices by probing them?
- How can we identify and stop attacks?
- Can we improve performance by analyzing traffic?
- Can users be tracked via their network traffic?
- Recently – Can machine learning techniques solve these problems?

Classic ML Pipeline


**Hypothesize
Problem**





**Gather
Traffic**




**Data
Processing**



**Engineer
Features**



**Train
Models**




Bespoke Solutions

Application Identification


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models




Bespoke Solutions

Application Identification


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models




Anomaly Detection


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models




Generalizable Solutions?


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models



Where We Are Headed


- Introduce nPrint, a generalizable packet representation that works across multiple traffic analysis tasks
- Combine nPrint and AutoML to create nPrintML, an open source system that generates full analysis pipelines

Classic ML Pipeline


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models




Features Are Expensive


Hypothesize
Problem





Gather
Traffic




Data
Processing



Engineer
Features



Train
Models




Are We Working Too Hard?


Hypothesize
Problem





Gather
Traffic




Data
Processing




Generic
Features



Train
Models



Goal Pipeline



Inspiration

- Image recognition
- Website fingerprinting on Tor traffic ^[1,2]

Network Traffic Issues

- Image recognition
- Website fingerprinting on Tor traffic [1,2]
- Problem – outside of Tor, network traffic not as simple!

Different Protocols

TCP Segment Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

UDP Datagram Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Length				Header and Data Checksum			

[3]

Different Values

TCP Segment Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

UDP Datagram Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Length				Header and Data Checksum			

Different Lengths

TCP Segment Header Format								
Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

UDP Datagram Header Format							
Bit #	0	7	8	15	16	23	24 31
0	Source Port				Destination Port		
32	Length				Header and Data Checksum		

Generic Packet Representation

The Semantic View

- Encode each header field as a feature

Semantic Representation: (IP / TCP) Packet					
IP Verison	IP IHL	IP ...	TCP Source Port	TCP ...	Payload
4	5	...	80	...	?

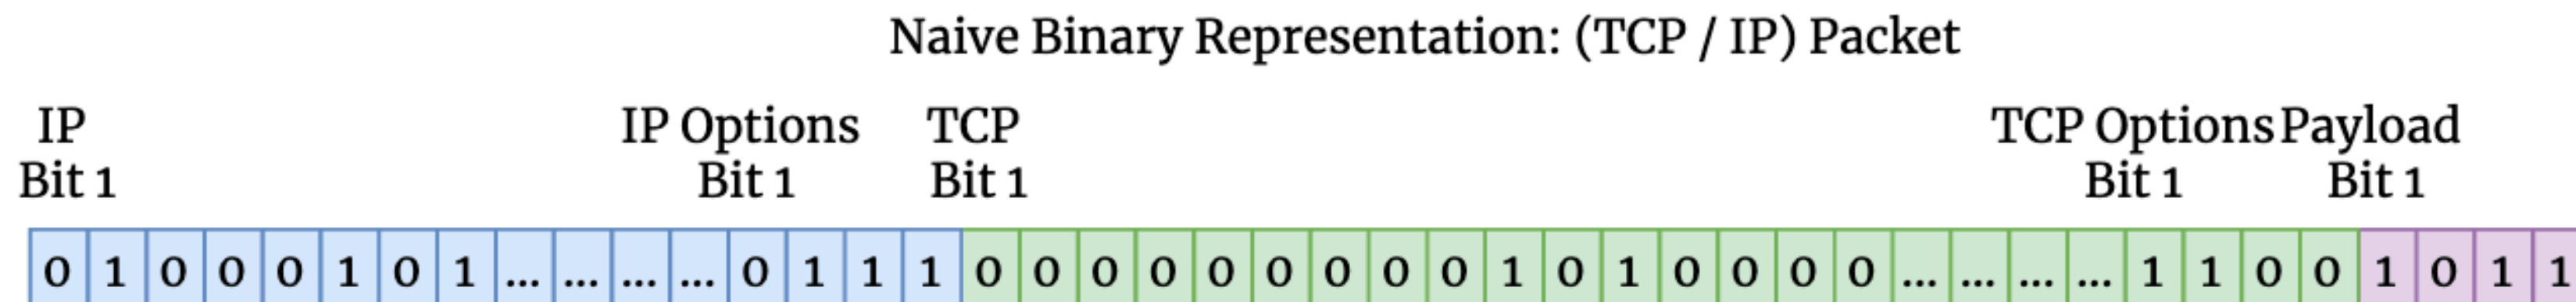

Semantic Issues

- Semi-structured fields
- Domain expertise
- Normalization
- Payloads

Semantic Representation: (IP / TCP) Packet					
IP Verison	IP IHL	IP ...	TCP Source Port	TCP ...	Payload
4	5	...	80	...	?


The Binary View

- Insight: packets are a collection of bits




Naive Noise

Naive Binary Representation: (TCP / IP) Packet




Naive Binary Representation: (TCP / IP) Packet: No Options




Naive Noise


Naive Binary Representation: (TCP / IP) Packet



Naive Binary Representation: (TCP / IP) Packet: No Options



Naive Binary Representation: (UDP / IP) Packet



nPrint

IPv4 480 Features	TCP 480 Features	UDP 64 Features	ICMP 64 Features	Payload <i>n</i> Features
Maximum Size of IPv4 Header (60 Bytes)	Maximum Size of TCP Header (60 Bytes)	Size of UDP Header (8 Bytes)	Size of ICMP Header (8 Bytes)	User Defined Number of Bytes

nPrint (TCP / IP) Packet

0	1	0	0	0	1	0	1	1	1	0	1	0	1	0	1	0	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	1	1	0
---	---	---	---	---	---	---	---	---	---	-----	-----	-----	-----	---	---	---	---	---	---	---	---	-----	-----	-----	-----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	---	---	---	-----	-----	-----	-----

nPrint (UDP / IP) Packet


0	1	0	0	0	1	0	1	1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	0	1	1	-1	-1	-1	-1	-1	-1	-1	-1	1	1	0
---	---	---	---	---	---	---	---	---	---	-----	-----	-----	-----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	---	---	---	-----	-----	-----	-----	----	----	----	----	----	----	----	----	---	---	---	-----	-----	-----	-----

Classic ML Pipeline


**Hypothesize
Problem**





**Gather
Traffic**




**Data
Processing**



**Engineer
Features**



**Train
Models**




nPrint Replaces Feature Engineering


Hypothesize
Problem





Gather
Traffic




Data
Processing



nPrint



Train
Models




New Bottleneck


Hypothesize
Problem





Gather
Traffic




Data
Processing



nPrint



Train
Models




Can We Automate This Step?


Hypothesize
Problem




Gather
Traffic




Data
Processing



nPrint



Train
Models



Classic Model Training

- Pick favorite model(s)
- Search some hyper-parameters for that model
- Choose best model


Automated Machine Learning

- Model selection
- Feature selection
- Hyperparameter search


AutoGluon AutoML

- Model ensembling achieves higher performance than other AutoML techniques^[8]
- Train models from 7 base classes
 - Random forests
 - DNN
 - KNN

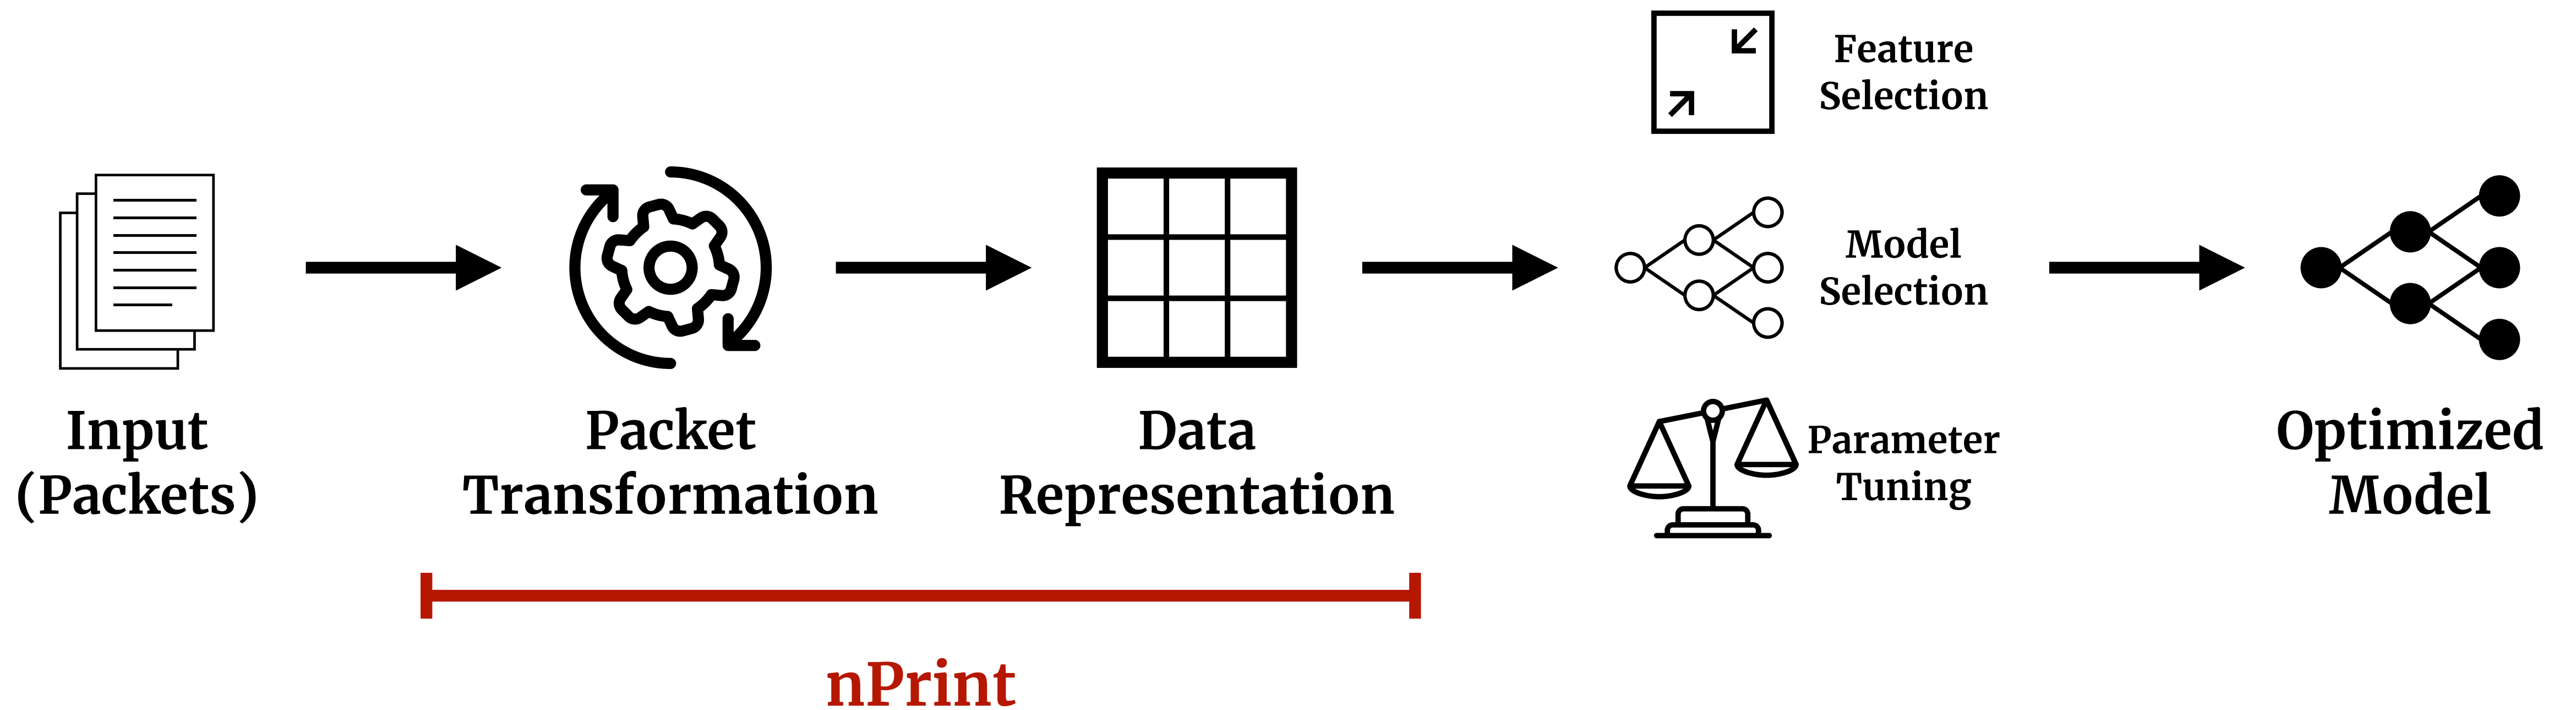
Original Goal Pipeline



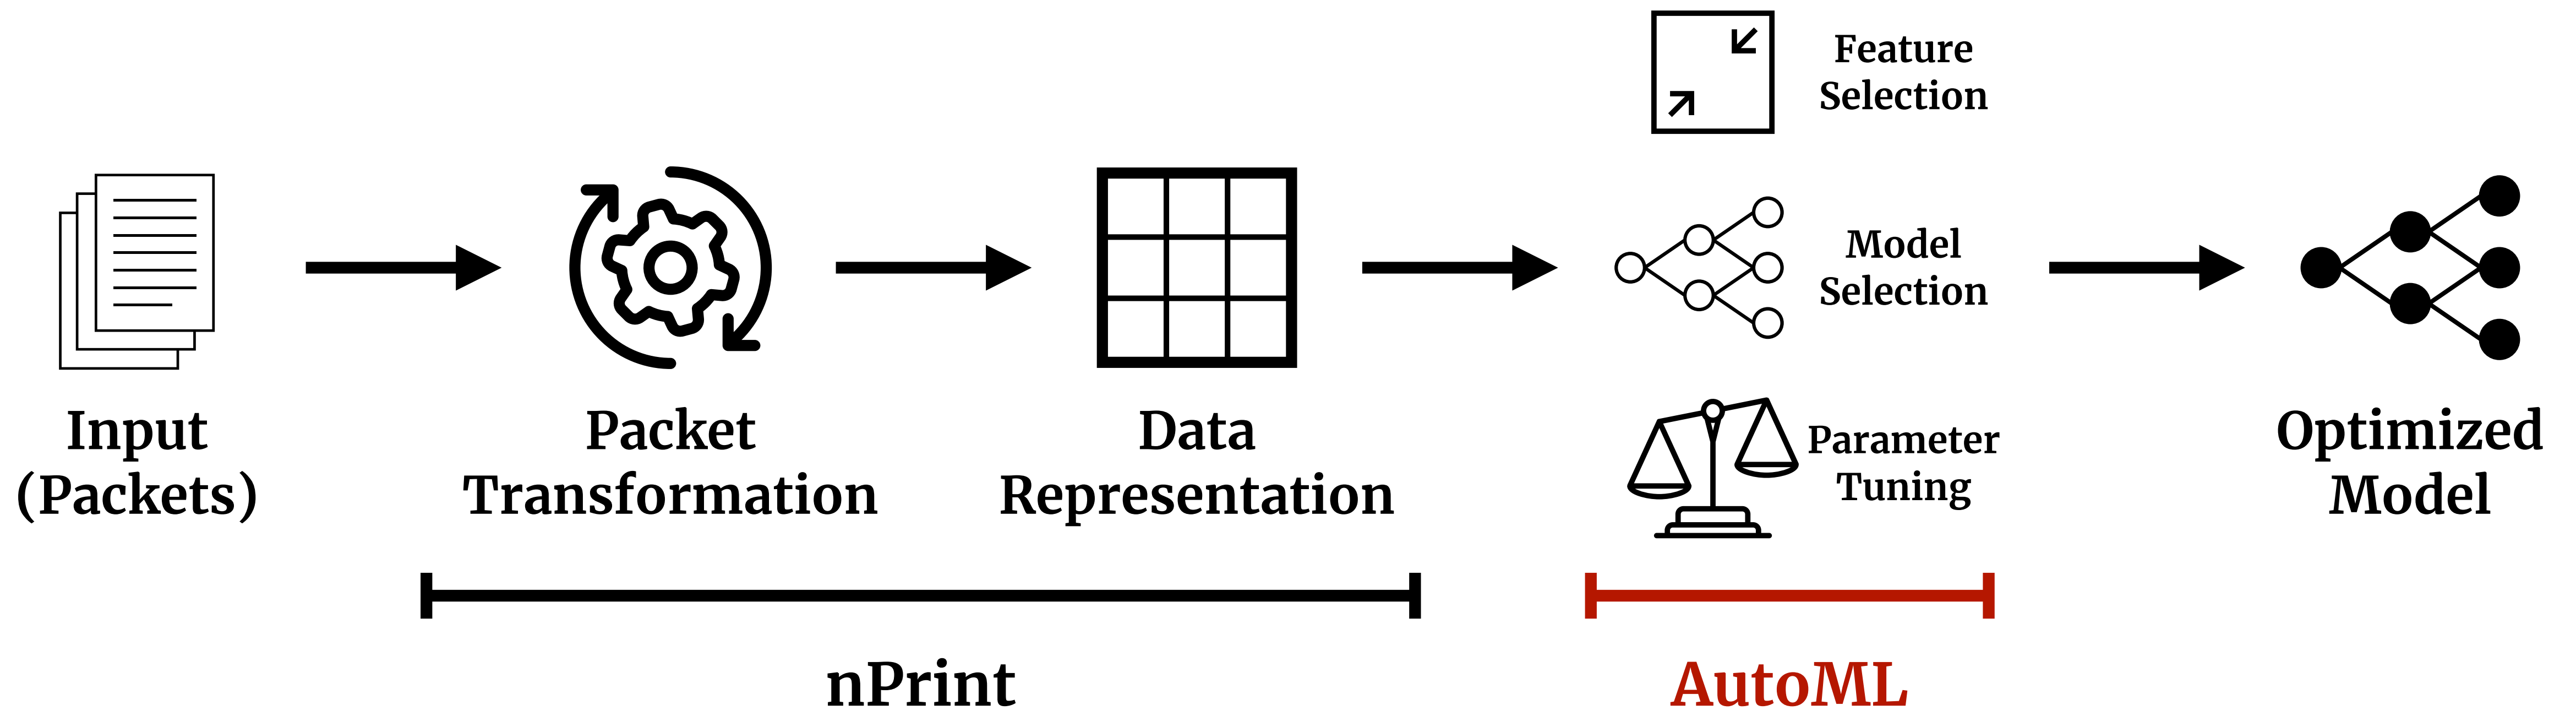
Detailed Traffic Analysis



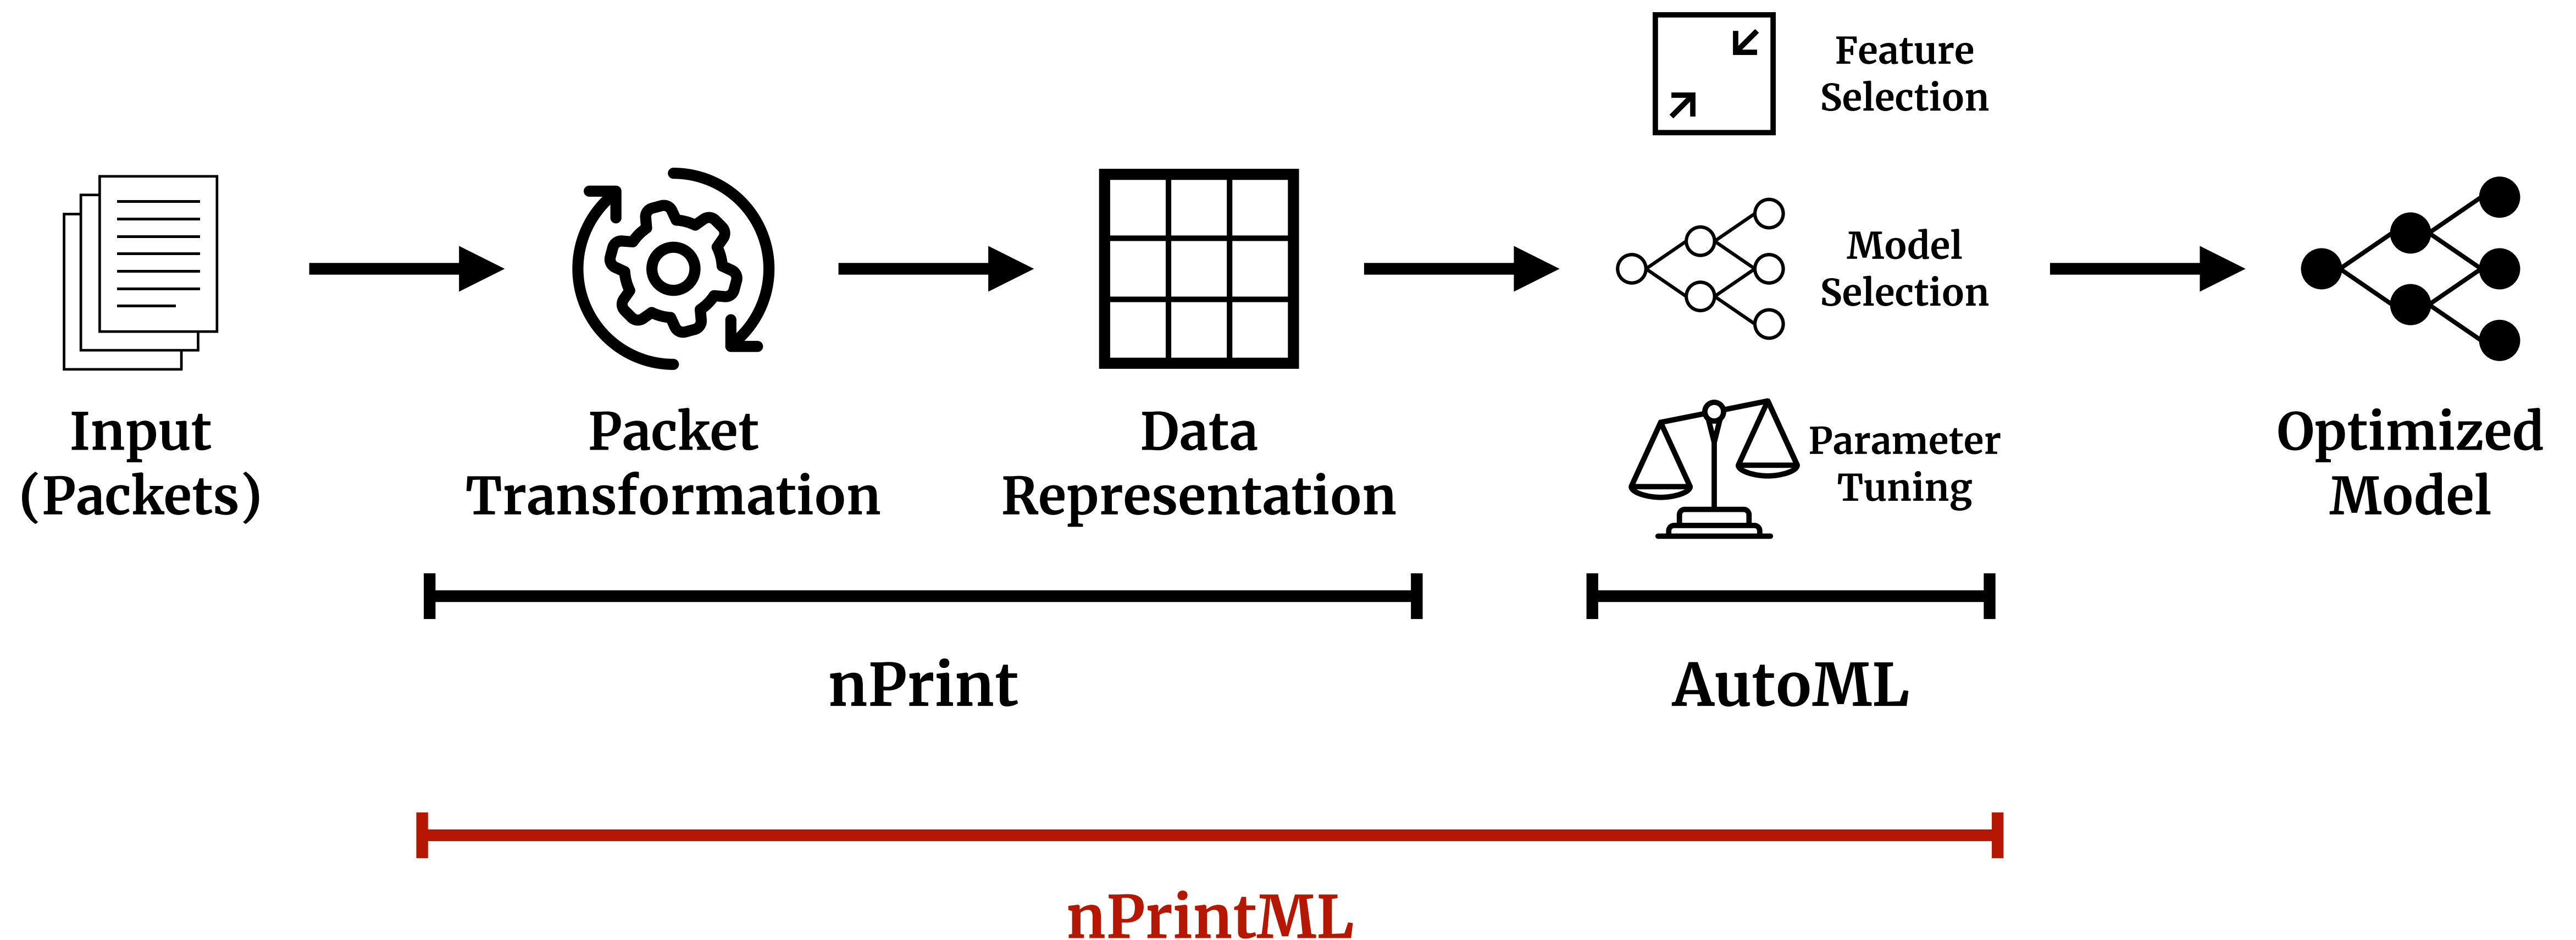
nPrint Transforms And Represents Packets



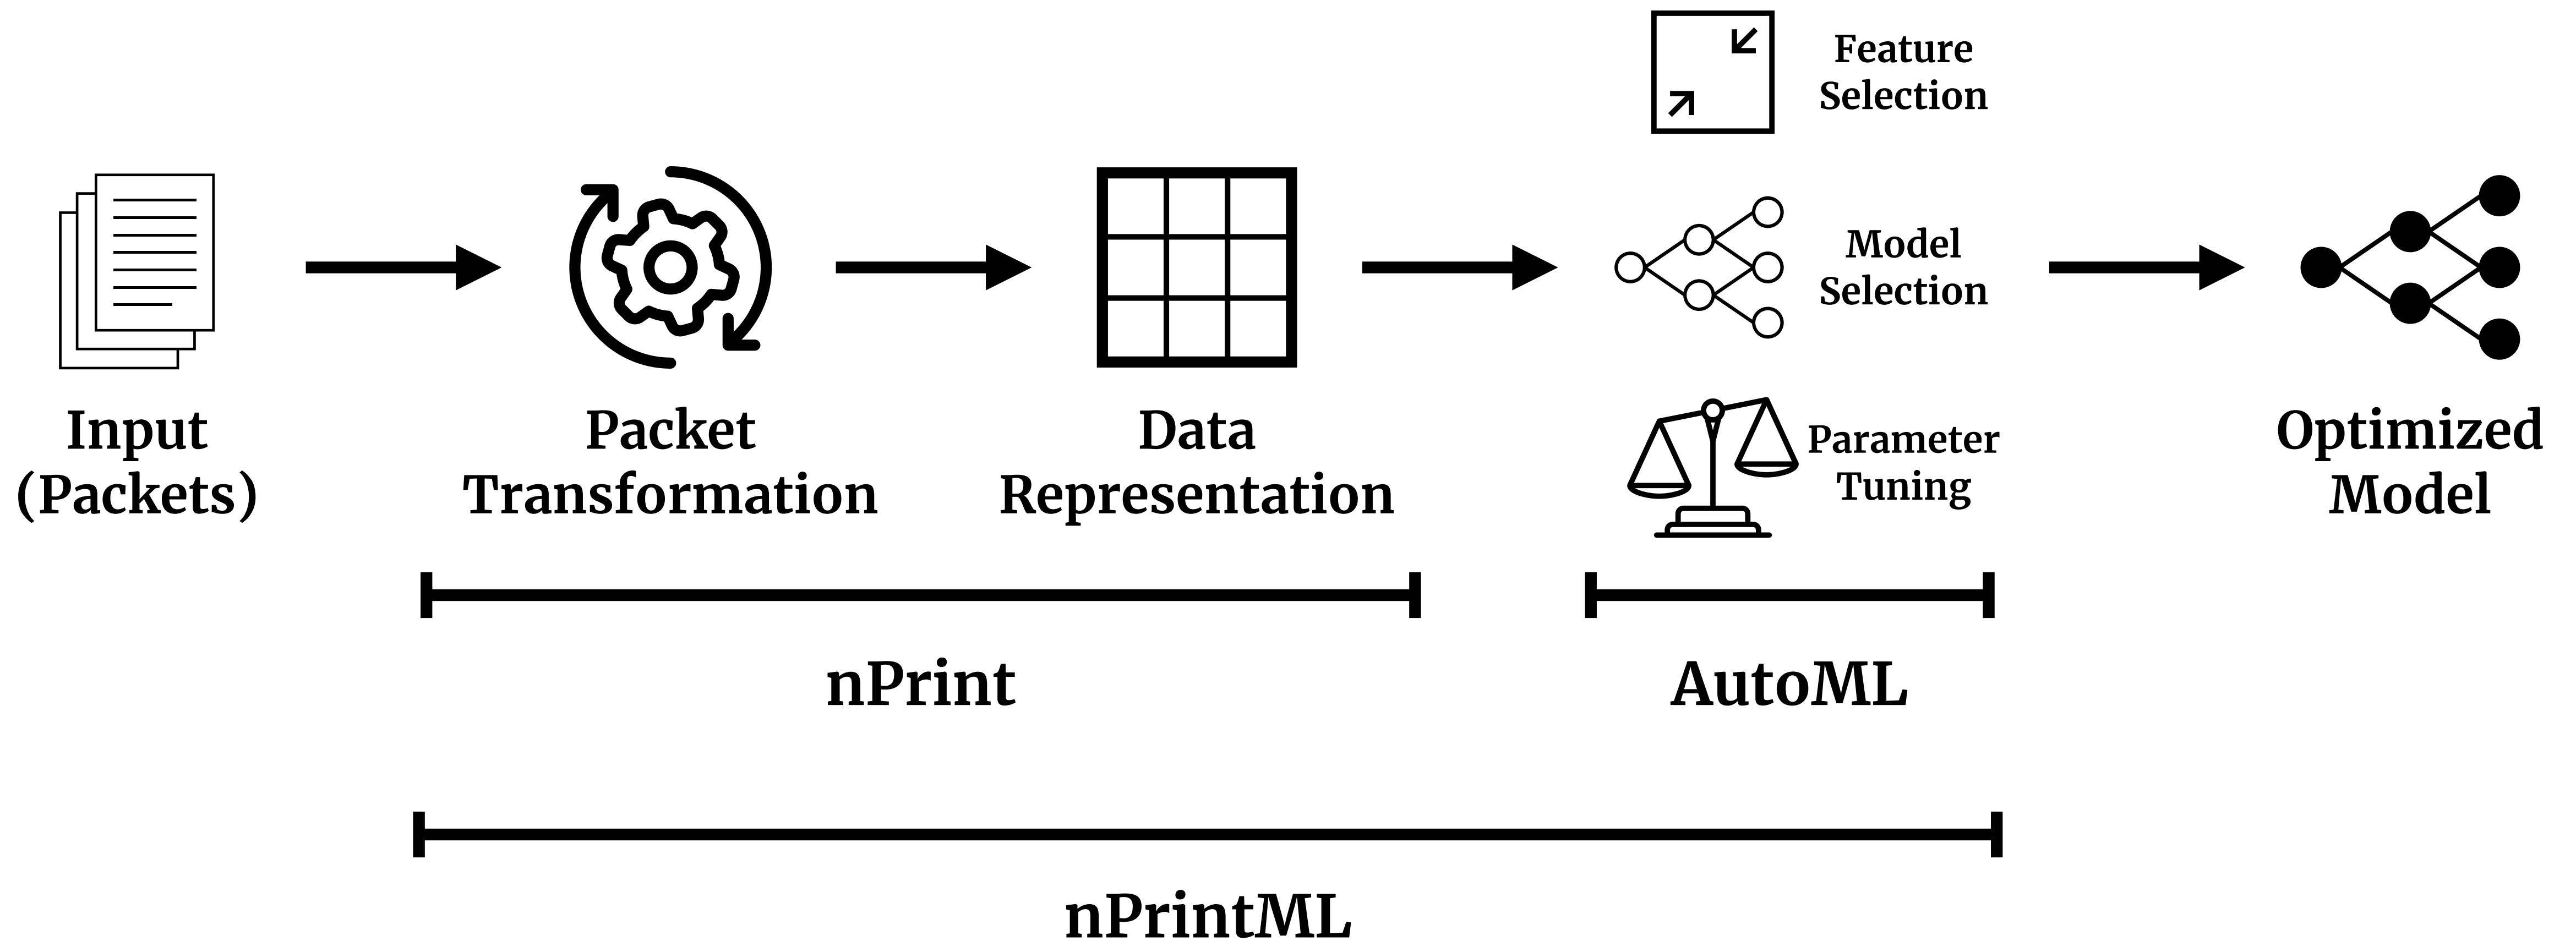
AutoML Finds The Best Model



nPrintML Combines Both!



Let's Try it!



Defining The Problem

- Remote device fingerprinting

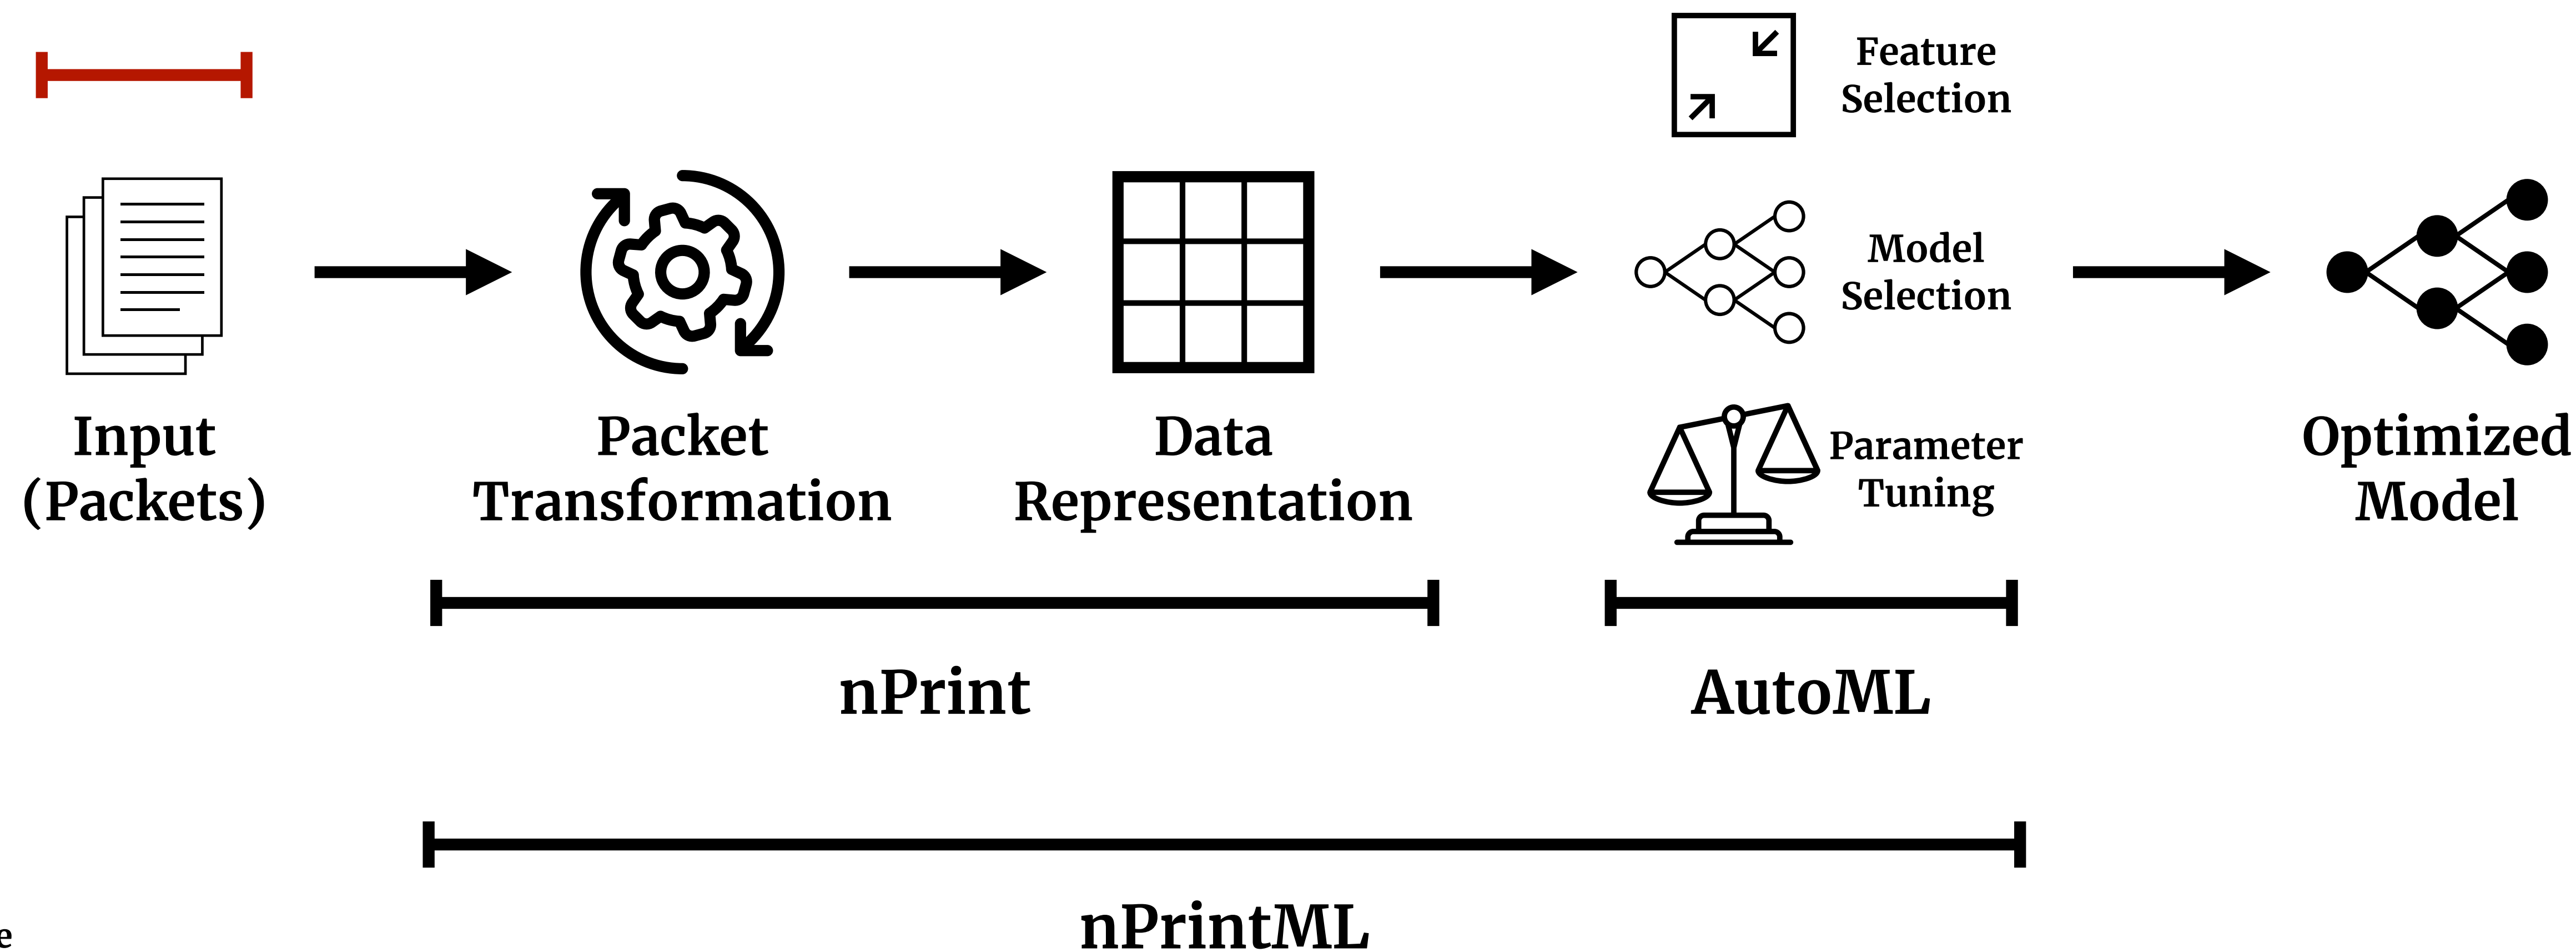
Building A Dataset

- Remote device fingerprinting
- Labeled Targets
 - Routers_[4]
 - IoT Devices (Shodan)


Labeled Dataset

Vendor	Device Type	Labeled Devices
Adtran	Network Device	1,449
Avtech	IoT Camera	2,152
Axis	IoT Camera	2,653
Chromecast	IoT Streaming	2,872
Cisco	Network Device	1,451
Dell	Network Device	1,449
H3C	Network Device	1,380
Huawei	Network Device	1,409
Juniper	Network Device	1,445
Lancom	Network Device	1,426
Mikrotik	Network Device	1,358
NEC	Network Device	1,450
Roku	IoT Streaming	2,403
Ubiquiti	Network Device	1,476
ZTE	Network Device	1,425


Gathering Traffic




Leveraging Nmap



Send 16 probes
(13 TCP), (2 ICMP), (1 UDP)




Extract hand
engineered fingerprint




Compare to database
using heuristic

Leveraging Nmap



Send 16 probes
(13 TCP), (2 ICMP), (1 UDP)



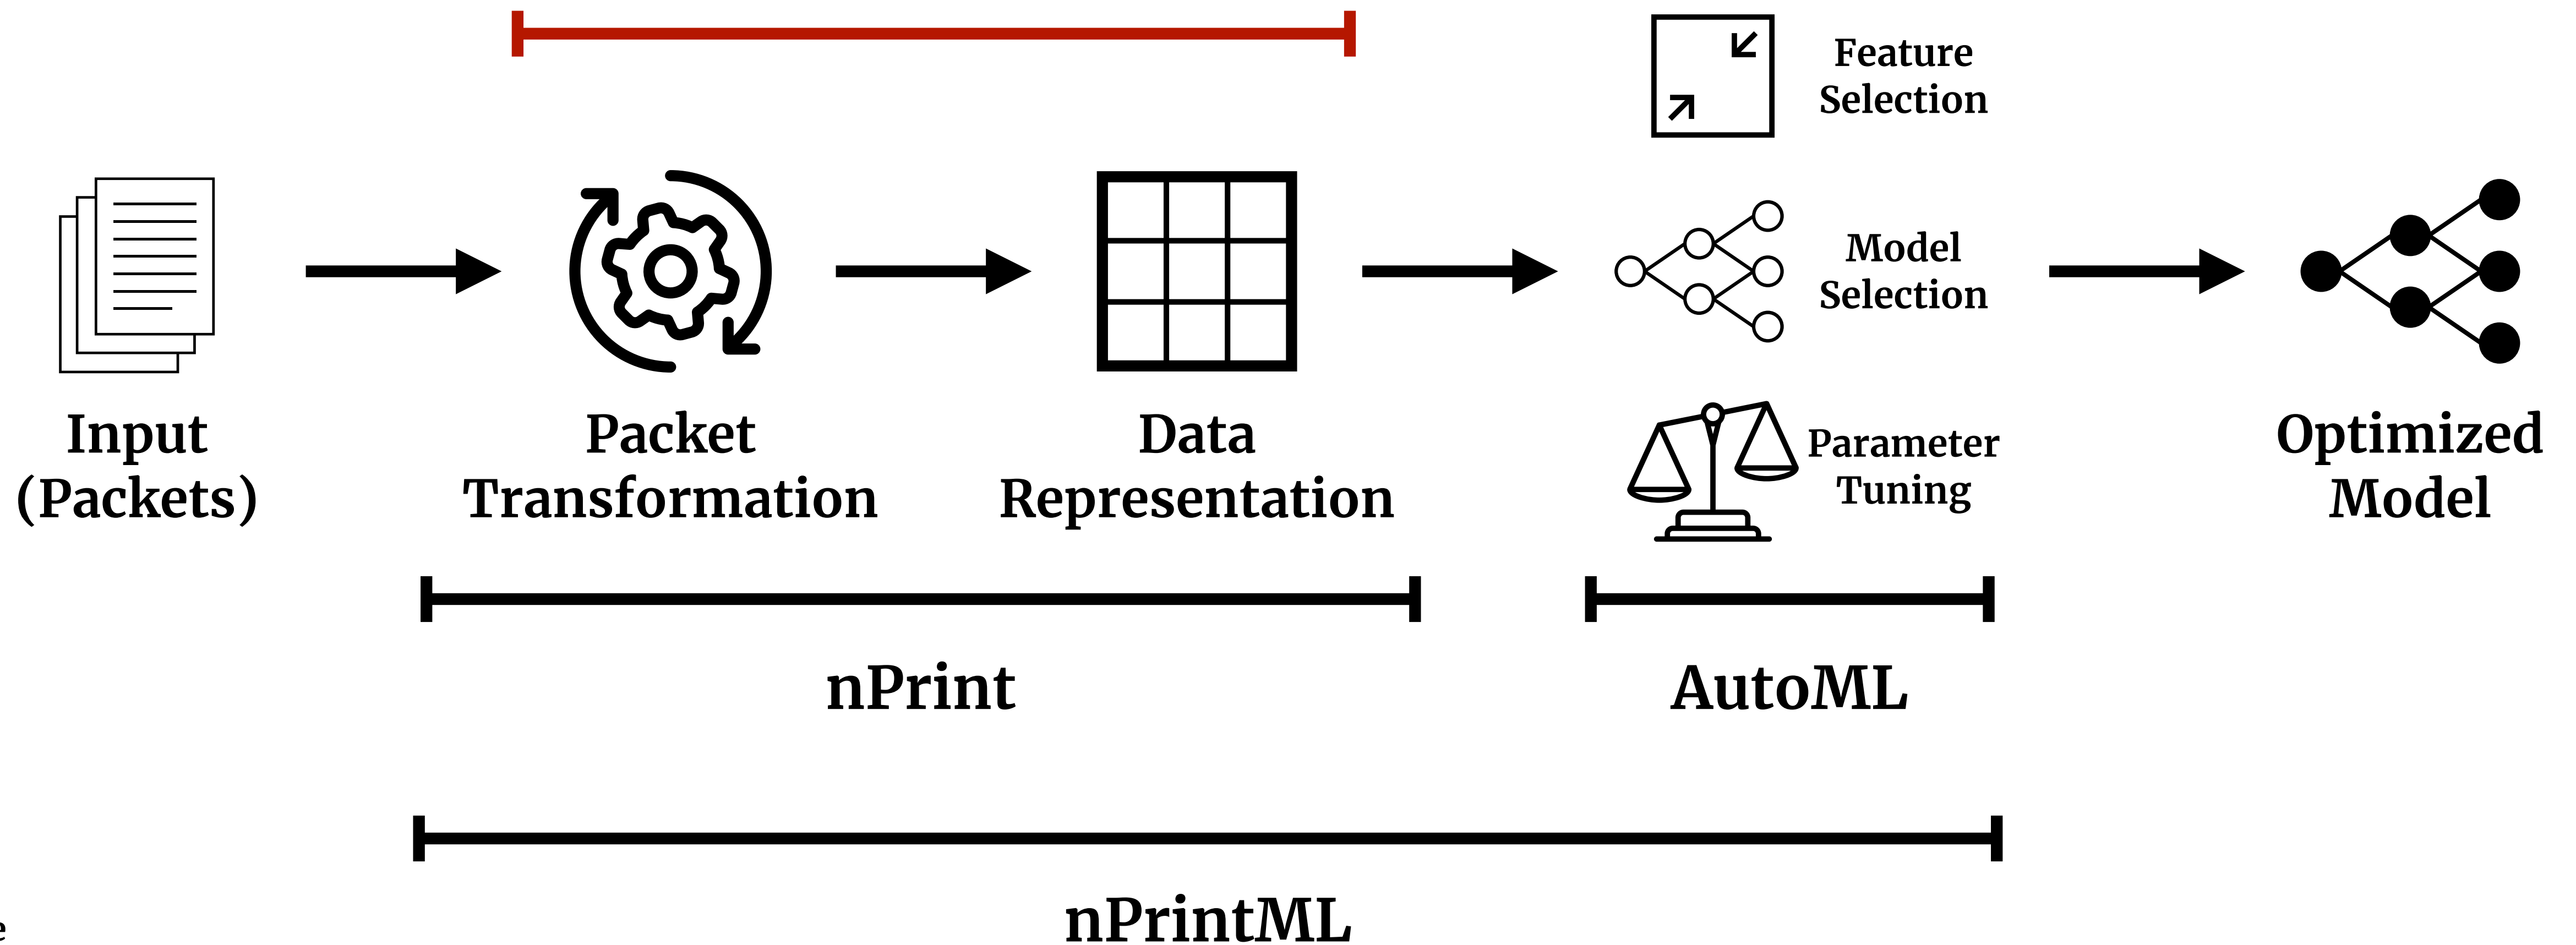
Extract hand
engineered fingerprint



Compare to database
using heuristic



Transforming Packets



nPrint Packet Transformation

- 21 uniquely named responses
- Sort responses by name and concatenate individual nPrints

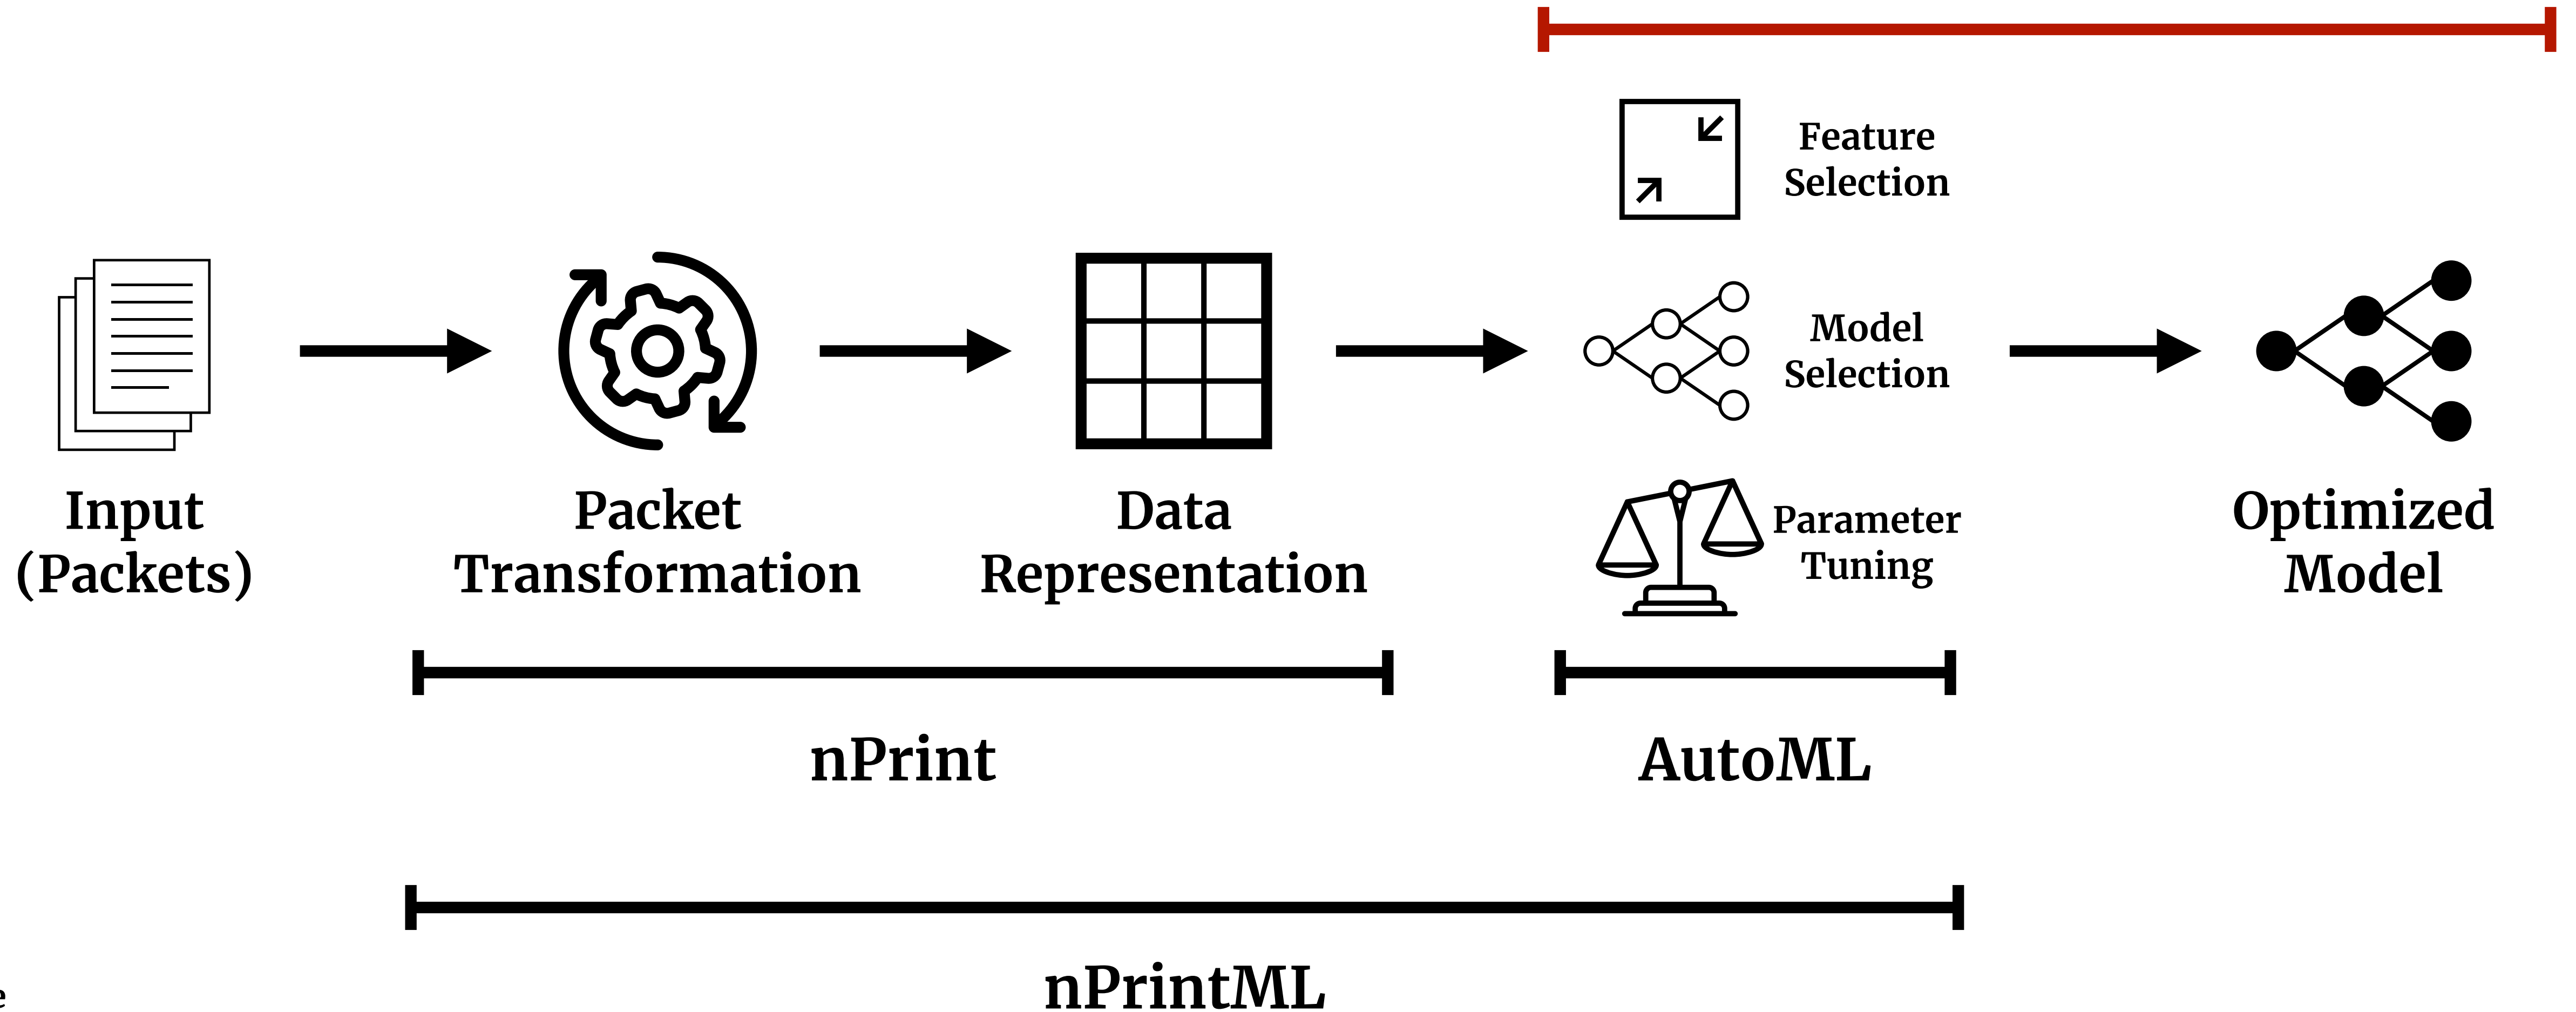
UDP Response
ICMP Response 1
ICMP Response 2
TCP Response 1
Response ...

21 Rows

Nmap Packet Transformation

Test Name	Summary	Nmap Weight
Explicit Congestion Notification	TCP Explicit Congestion control flag.	100
ICMP Response Code	ICMP Response Code.	100
Integrity of returned probe IP Checksum	Valid checksum in an ICMP port unreachable.	100
Integrity of returned probe UDP Checksum	UDP header checksum received match.	100
IP ID Sequence Generation Algorithm	Algorithm for IP ID.	100
IP Total Length	Total length of packet.	100
Responsiveness	Target responded to a given probe.	100
Returned probe IP ID value	IP ID value.	100
Returned Probe IP Total Length	IP Length of an ICMP port unreachable.	100
TCP Timestamp Option Algorithm	TCP timestamp option algorithm.	100
Unused Port unreachable Field Nonzero	Last 4 bytes of ICMP port unreachable message not zero.	100
Shared IP ID Sequence Boolean	Shared IP ID Sequence between TCP and ICMP.	80
TCP ISN Greatest Common Divisor	Smallest TCP ISN increment.	75
Don't Fragment ICMP	IP Don't Fragment bit for ICMP probes.	40
TCP Flags	TCP flags.	30
TCP ISN Counter Rate	Average rate of increase for the TCP ISN.	25
TCP ISN Sequence Predictability Index	Variability in the TCP ISN.	25
IP Don't Fragment Bit	IP Don't Fragment bit.	20
TCP Acknowledgment Number	TCP acknowledgment number.	20
TCP Miscellaneous Quirks	TCP implementations, e.g, reserved field in TCP header.	20
TCP Options Test	TCP header options, preserving order.	20
TCP Reset Data Checksum	Checksum of data in TCP reset packet.	20
TCP Sequence Number	TCP sequence number.	20
IP Initial Time-To-Live	IP initial time-to-live.	15
TCP Initial Window Size	TCP window size.	15

Training Models

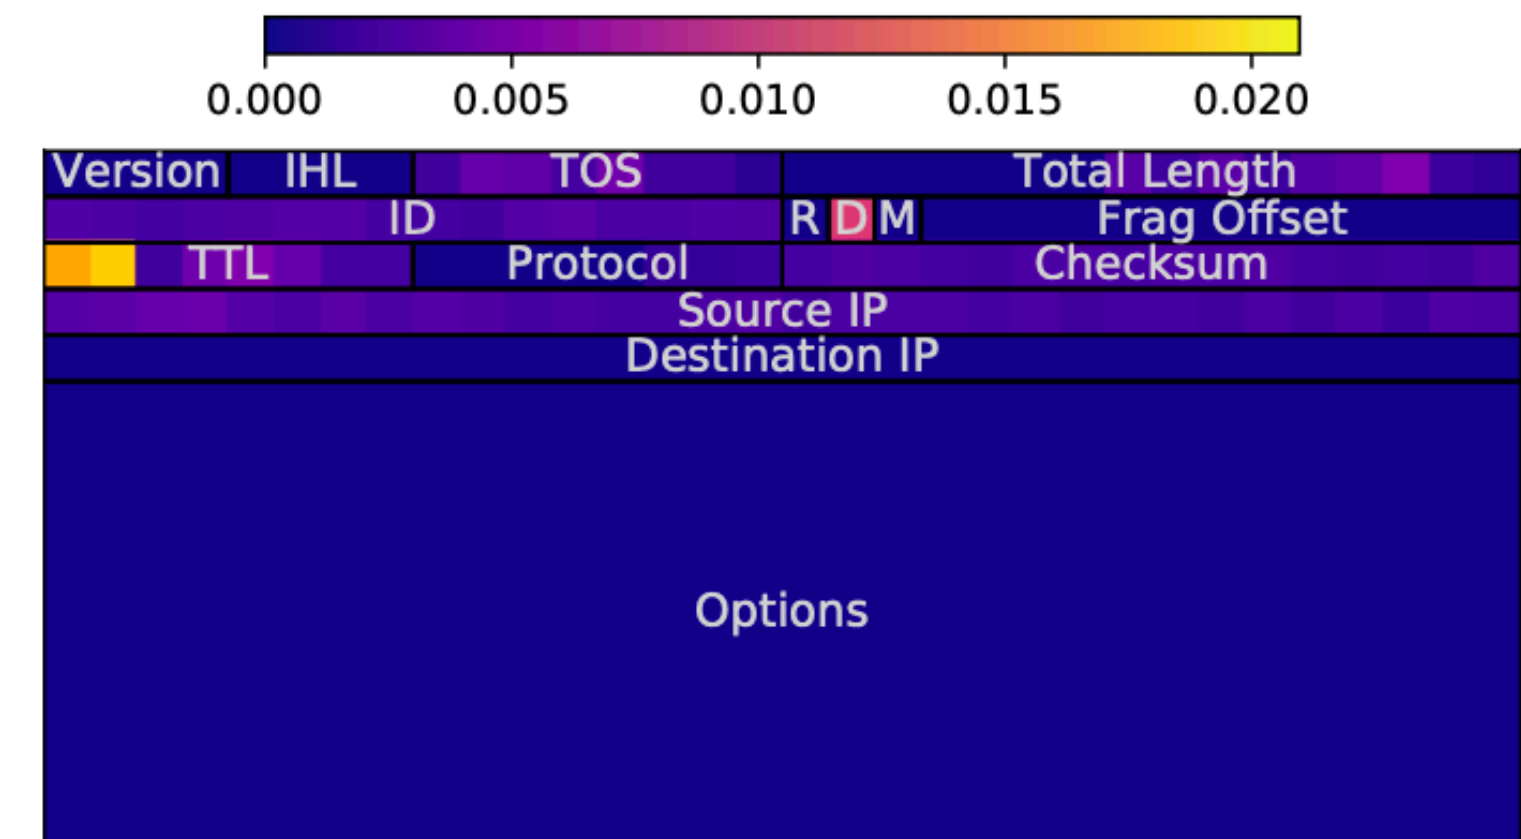


nPrint Outperforms Nmap

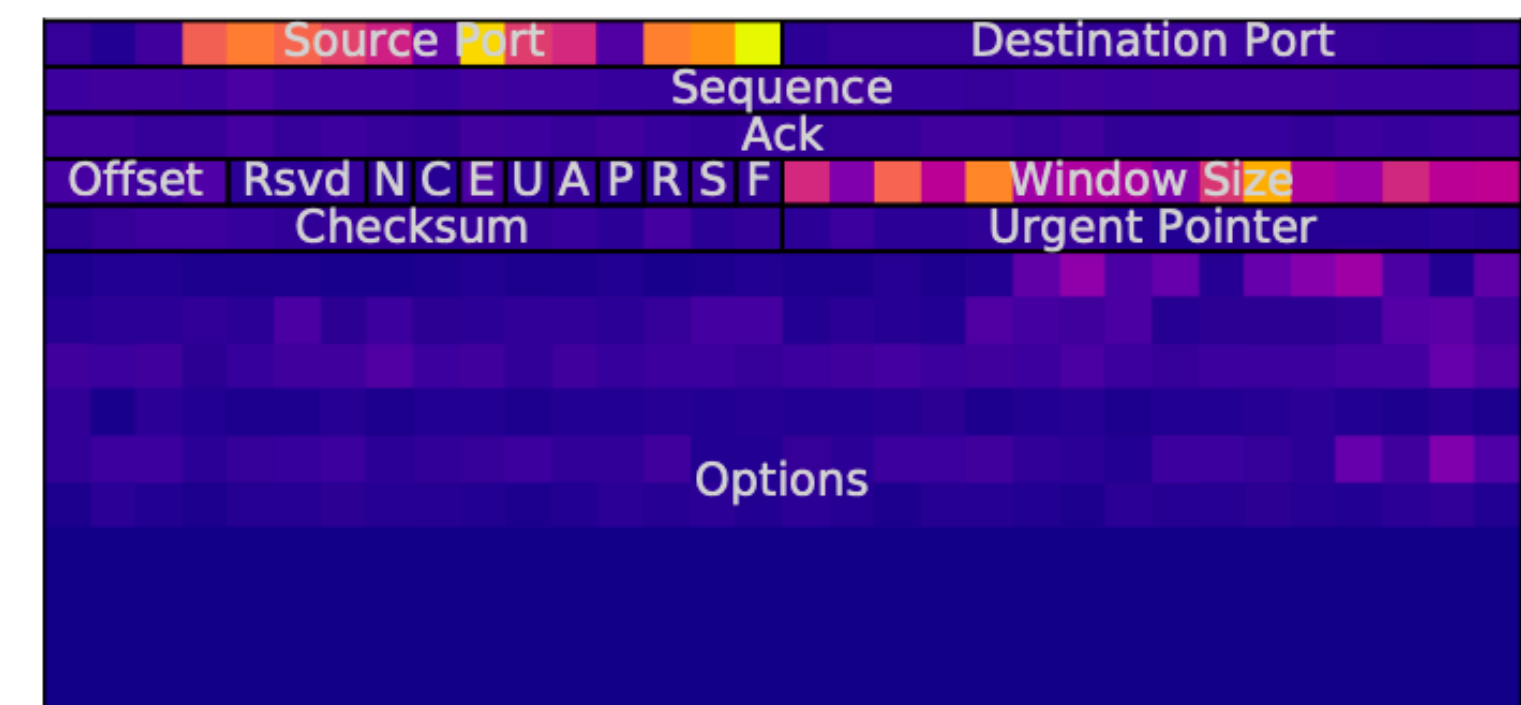
Representation	Balanced Accuracy	ROC AUC	F1
nPrint	95.4	99.7	95.5
Nmap	92.7	99.3	92.9

nPrint Enables Interpretable Machine Learning

- Map features to packet header semantics!
- Automatically learn
 - IP TTL
 - TCP options, window size
 - Source port identifies IoT vs Routers



(a) IPv4



(b) TCP



(c) ICMP

nPrintML’s Breadth

Problem Overview			nPrintML					Comparison	
Description	Dataset	# Classes	Configuration eAppendix A.4)	Sample Size (# Packets)	Balanced Accuracy	ROC AUC	Macro F1	Score	Source
Active Device Fingerprinting (§5.1)	Network Device Dataset [22]	15	-4 -t -i	21	95.4	99.7	95.5	92.9 (Macro-F1)	ML-Enhanced Nmap [31]

8 Discrete Case Studies

Problem Overview			nPrintML					Comparison		
Description	Dataset	# Classes	Configuration eAppendix A.4)	Sample Size (# Packets)	Balanced Accuracy	ROC AUC	Macro F1	Score	Source	
Active Device Fingerprinting (§5.1)	Network Device Dataset [22]	15	-4 -t -i	21	95.4	99.7	95.5	92.9 (Macro-F1)	ML-Enhanced Nmap [31]	
Passive OS Detection (§5.2)	CICIDS 2017 [48]	3	-4 -t	1	99.5	99.9	99.5	81.3 (Macro-F1)	p0f [40]	
		10		99.9	100	99.9				
		13		100	99.9	100	99.9	No Previous Work		
Application Identification via DTLS Handshakes (§5.3)	DTLS Handshakes [32]	7	-4	43	99.8	96.9	99.7	99.8 (Average Accuracy)	Hand-Curated Features [32]	
			-u		99.9	99.7	99.5			
			-p 10		95.0	78.8	77.4			
			-p 25		99.9	99.7	99.7			
			-p 100		99.9	99.7	99.7			
-4 -u -p 10	99.8	99.9	99.8							
Malware Detection for IoT Traces (§5.4.1)	netML IoT [6, 28]	2	-4 -t -u	10	92.4	99.5	93.2	99.9 (True Positive Rate) 39.7 (Balanced F1)	NetML Challenge Leaderboard [37]	
		19			86.1	96.9	84.1			
Type of Traffic in Capture (§5.4.1)	netML Non-VPN [6, 12]	7	-4 -t -u -p 10	10	81.9	98.0	79.5	67.3 (Balanced F1)		
		18			76.1	94.2	75.8	42.1 (Balanced F1)		
		31			66.2	91.3	63.7	34.9 (Balanced F1)		
		31			60.9	92.2	57.6	34.9 (Balanced F1)		
Intrusion Detection (§5.4.1)	netML CICIDS 2017 [6, 48]	2	-4 -t -u	5	99.9	99.9	99.9	98.9 (True Positive Rate) 99.2 (Balanced F1)		
		8			99.9	99.9	99.9			
Determine Country of Origin for Android & iOS Application Traces (§5.4.2)	Cross Platform [44]	3	-4 -t -u -p 50	25	96.8	90.2	90.4	No Previous Work		
Identify streaming video (DASH) service via device SYN packets (§5.4.3)	Streaming Video Providers [10]	4	-4 -t -u -R	10	77.9	96.0	78.9	No Previous Work		
				25	90.2	98.6	90.4			
				50	98.4	99.9	98.6			

Outperforming hand-engineered solutions

Problem Overview			nPrintML					Comparison	
Description	Dataset	# Classes	Configuration eAppendix A.4)	Sample Size (# Packets)	Balanced Accuracy	ROC AUC	Macro F1	Score	Source
Active Device Fingerprinting (§5.1)	Network Device Dataset [22]	15	-4 -t -i	21	95.4	99.7	95.5	92.9 (Macro-F1)	ML-Enhanced Nmap [31]
Passive OS Detection (§5.2)	CICIDS 2017 [48]	3	-4 -t	1	99.5	99.9	99.5	81.3 (Macro-F1)	p0f [40]
				10	99.9	100	99.9		
		13		100	99.9	100	99.9		
Application Identification via DTLS Handshakes (§5.3)	DTLS Handshakes [32]	7	-4 -u -p 10 -p 25 -p 100 -4 -u -p 10	43	77.1	97.5	76.9	99.8 (Average Accuracy)	Hand-Curated Features [32]
					99.8	96.9	99.7		
					99.9	99.7	99.5		
					95.0	78.8	77.4		
					99.9	99.7	99.7		
Malware Detection for IoT Traces (§5.4.1)	netML IoT [6, 28]	2	-4 -t -u	10	99.9	99.9	99.9	98.9 (True Positive Rate)	NetML Challenge Leaderboard [37]
		19			99.9	99.9	99.9	99.2 (Balanced F1)	
Type of Traffic in Capture (§5.4.1)	netML Non-VPN [6, 12]	7	-4 -t -u -p 10	10	81.9	98.0	79.5	67.3 (Balanced F1)	
					76.1	94.2	75.8	42.1 (Balanced F1)	
		18			66.2	91.3	63.7	34.9 (Balanced F1)	
		31			60.9	92.2	57.6	34.9 (Balanced F1)	
Intrusion Detection (§5.4.1)	netML CICIDS 2017 [6, 48]	2	-4 -t -u	5	99.9	99.9	99.9	98.9 (True Positive Rate)	
		8			99.9	99.9	99.9	99.2 (Balanced F1)	
Determine Country of Origin for Android & iOS Application Traces (§5.4.2)	Cross Platform [44]	3	-4 -t -u -p 50	25	96.8	90.2	90.4	No Previous Work	
Identify streaming video (DASH) service via device SYN packets (§5.4.3)	Streaming Video Providers [10]	4	-4 -t -u -R	10	77.9	96.0	78.9	No Previous Work	
				25	90.2	98.6	90.4		
				50	98.4	99.9	98.6		

nPrint Is Open Source

- 8 protocols implemented
- Relative & absolute timestamps
- Input formats – live capture, PCAP, scan data, nPrints

nPrintML Is Open Source

- Application Identification
 - nprintml —pcap-dir pcaps/ -L labels.csv -a pcap -4 -u -p 10
- Passive OS detection
 - nprintml -P traffic.pcap -L labels.csv -a index -4 -t

Thank You!

github.com/nprint/

References

1. Rimmer, Vera, et al. "Automated website fingerprinting through deep learning." *25th Annual Network & Distributed System Security Symposium*. NDSS, 2018
2. Oh, Se Eun, Saikrishna Sunkam, and Nicholas Hopper. "p1-FP: Extraction, Classification, and Prediction of Website Fingerprints with Deep Learning." *Proceedings on Privacy Enhancing Technologies* 2019.3 (2019): 191–209.
3. <https://skminhaj.wordpress.com/2016/02/15/tcp-segment-vs-udp-datagram-header-format/>
4. <https://arxiv.org/pdf/2006.13086.pdf>
5. <https://www.shodan.io/>
6. Iman Sharafaldin, Arash Habibi Lashkari, and Ali A. Ghorbani, "Toward Generating a New Intrusion Detection Dataset and Intrusion Traffic Characterization", 4th International Conference on Information Systems Security and Privacy (ICISSP), Portugal, January 2018
7. Sambasivan, Nithya, et al. "'Everyone wants to do the model work, not the data work': Data Cascades in High-Stakes AI." *proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. 2021.
8. Erickson, Nick, et al. "Autogluon-tabular: Robust and accurate automl for structured data." *arXiv preprint arXiv:2003.06505* (2020).